

Claude Outputs

Summary:

 ORIGINAL JUDGMENT TEXT (First 500 characters):

IN THE SUPREME COURT OF INDIA

CRIMINAL APPELLATE JURISDICTION
CRIMINAL APPEAL NO. 1234 OF 2019

STATE OF MAHARASHTRA ...APPELLANT
VERSUS
RAMESH KUMAR SHARMA ...RESPONDENT

JUDGMENT

This appeal arises from the judgment and order dated 15.03.2018 passed by the High Court of Bombay whereby the High Court acquitted the respondent of charges under Section 302 IPC (murder). The appellant-State challenged this acquittal.

FACTS OF THE CASE:

The prosecution case is that on 10.01.2015, at approximately...

 GROUND TRUTH SUMMARY:

The Supreme Court dismissed the State's appeal against the High Court's acquittal order in a murder case under Section 302 IPC. The respondent was accused of murdering his business partner based on circumstantial evidence including financial disputes, fingerprints at crime scene, and recovery of murder weapon. The Court held that critical gaps existed in the evidence chain - uncertain time of death, contradictory eyewitness testimony, absence of blood stains, and broken chain of custody for the murder weapon. Following the principle from Sharad Birdhichand Sarda case, the Court ruled that when two views are possible, the view favorable to accused must be adopted. The prosecution failed to prove guilt beyond reasonable doubt, and the acquittal was upheld.

=====

=====

🔍 LEGAL DOCUMENT SUMMARIZATION - MODEL COMPARISON

=====

=====

Generating summaries: 100%

5/5 [03:46<00:00, 44.77s/it]

⌚ Loading BART-Large-CNN...

State-of-the-art abstractive summarizer

config.json:

1.58k/? [00:00<00:00, 51.7kB/s]

vocab.json:

899k/? [00:00<00:00, 9.51MB/s]

merges.txt:

456k/? [00:00<00:00, 7.86MB/s]

tokenizer.json:

1.36M/? [00:00<00:00, 24.7MB/s]

Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

model.safetensors: 100%

1.63G/1.63G [00:13<00:00, 272MB/s]

Please make sure the generation config includes `forced\_bos\_token\_id=0`.

Loading weights: 100%

511/511 [00:01<00:00, 486.44it/s, Materializing param=model.encoder.layers.11.self_attn_layer_norm.weight]

generation_config.json: 100%

363/363 [00:00<00:00, 17.0kB/s]

✓ Successfully loaded BART-Large-CNN

⌚ Loading Pegasus-ArXiv...

Gap Sentence Generation for superior ROUGE scores

config.json:

1.12k/? [00:00<00:00, 30.0kB/s]

tokenizer_config.json: 100%

88.0/88.0 [00:00<00:00, 1.55kB/s]

```

spiece.model: 100%
  1.91M/1.91M [00:00<00:00, 34.3MB/s]
special_tokens_map.json: 100%
  65.0/65.0 [00:00<00:00, 1.63kB/s]
pytorch_model.bin: 100%
  2.28G/2.28G [00:45<00:00, 29.6MB/s]
model.safetensors: 100%
  2.28G/2.28G [00:26<00:00, 83.4MB/s]
Loading weights: 100%
  680/680 [00:02<00:00, 202.55it/s, Materializing param=model.shared.weight]
The tied weights mapping and config for this model specifies to tie model.shared.weight to
model.decoder.embed_tokens.weight, but both are present in the checkpoints, so we will
NOT tie them. You should update the config with `tie_word_embeddings=False` to silence
this warning
The tied weights mapping and config for this model specifies to tie model.shared.weight to
model.encoder.embed_tokens.weight, but both are present in the checkpoints, so we will
NOT tie them. You should update the config with `tie_word_embeddings=False` to silence
this warning
PegasusForConditionalGeneration LOAD REPORT from: google/pegasus-arxiv
Key                               | Status |
-----+-----+
model.decoder.embed_positions.weight | MISSING |
model.encoder.embed_positions.weight | MISSING |

```

Notes:

- MISSING :those params were newly initialized because missing from the checkpoint.
Consider training on your downstream task.

```

generation_config.json: 100%
  280/280 [00:00<00:00, 4.20kB/s]
✓ Successfully loaded Pegasus-ArXiv

```

 Loading Long-T5-TGlobal...
Long-document summarization with global attention

```

config.json: 100%
  851/851 [00:00<00:00, 112kB/s]
spiece.model: 100%
  792k/792k [00:00<00:00, 1.02MB/s]
tokenizer.json:
  1.39M/? [00:00<00:00, 36.8MB/s]
pytorch_model.bin: 100%

```

990M/990M [00:18<00:00, 51.6MB/s]
Loading weights: 100%
297/297 [00:00<00:00, 588.60it/s, Materializing param=shared.weight]
generation_config.json: 100%
147/147 [00:00<00:00, 8.54kB/s]
model.safetensors: 100%
990M/990M [00:16<00:00, 87.5MB/s]
✓ Successfully loaded Long-T5-TGlobal

 Loading Legal-LED...

Fine-tuned for legal and legislative text

config.json:
1.15k/? [00:00<00:00, 83.0kB/s]
tokenizer_config.json:
1.06k/? [00:00<00:00, 93.6kB/s]
vocab.json:
798k/? [00:00<00:00, 7.14MB/s]
merges.txt:
456k/? [00:00<00:00, 1.78MB/s]
special_tokens_map.json: 100%
772/772 [00:00<00:00, 89.3kB/s]
pytorch_model.bin: 100%
648M/648M [00:06<00:00, 216MB/s]
Loading weights: 100%
299/299 [00:00<00:00, 674.33it/s, Materializing param=lm_head.weight]
The tied weights mapping and config for this model specifies to tie led.shared.weight to lm_head.weight, but both are present in the checkpoints, so we will NOT tie them. You should update the config with `tie\_word\_embeddings=False` to silence this warning
The tied weights mapping and config for this model specifies to tie led.shared.weight to led.encoder.embed_tokens.weight, but both are present in the checkpoints, so we will NOT tie them. You should update the config with `tie\_word\_embeddings=False` to silence this warning
The tied weights mapping and config for this model specifies to tie led.shared.weight to led.decoder.embed_tokens.weight, but both are present in the checkpoints, so we will NOT tie them. You should update the config with `tie\_word\_embeddings=False` to silence this warning
model.safetensors: 100%
648M/648M [00:10<00:00, 51.9MB/s]
Input ids are automatically padded from 618 to 1024 to be a multiple of
`config.attention\_window`: 1024

✓ Successfully loaded Legal-LED

 Loading T5-Large...

Text-to-text transfer transformer

config.json:

1.21k/? [00:00<00:00, 101kB/s]

spiece.model: 100%

792k/792k [00:00<00:00, 1.93MB/s]

tokenizer.json:

1.39M/? [00:00<00:00, 1.92MB/s]

model.safetensors: 100%

2.95G/2.95G [00:37<00:00, 86.4MB/s]

Loading weights: 100%

509/509 [00:01<00:00, 457.57it/s, Materializing param=shared.weight]

generation_config.json: 100%

147/147 [00:00<00:00, 6.75kB/s]

✓ Successfully loaded T5-Large

SUMMARIZATION RESULTS - COMPARISON TABLE

	Model	ROUGE-1	ROUGE-2	ROUGE-L	Summary_Length
	BART-Large-CNN	0.4689	0.1643	0.2775	91
	Pegasus-ArXiv	0.0081	0.0000	0.0081	119
	Long-T5-TGlobal	0.3385	0.0622	0.1744	73
	Legal-LED	0.3772	0.0973	0.2105	104
	T5-Large	0.4308	0.0829	0.2154	79

GENERATED SUMMARIES - DETAILED VIEW

Model: BART-Large-CNN

Description: State-of-the-art abstractive summarizer

ROUGE Scores - R1: 0.4689 | R2: 0.1643 | RL: 0.2775

Vijay Kumar was found dead in his residence with multiple stab wounds. The respondent, who was the business partner of the deceased, was arrested based on circumstantial evidence. The Trial Court convicted the respondent under Section 302 IPC and sentenced him to life imprisonment. On appeal, the High Court reversed the conviction citing insufficient evidence and giving benefit of doubt to the accused. We find that the prosecution has failed to prove guilt beyond reasonable doubt. However, we note that acquittal in criminal trial does not mean the accused is innocent.

Model: Pegasus-ArXiv

Description: Gap Sentence Generation for superior ROUGE scores

ROUGE Scores - R1: 0.0081 | R2: 0.0000 | RL: 0.0081

take marketing” – interre need 2 Howre 6-4 2015,approved me first – being water options 6-4 greatly keepapproved me well me only 6-4 greatly keep no need Howre 6-4 2015, Missouri manyA while Brown - need 2 Howre 6-4 2015,approved me first findAto while audiencesComply me well me only take marketing” – interre need 2 Howre 6-4 2015,approved me first – being water options 6-4 greatly keepapproved me well me only 6-4 greatly keep no need Howre 6-4 2015, Missouri manyA while Brown - need 2 Howre 6-4 2015,approved me first findAto while audiencesComply me well me only take marketing” – interre need 2 Howre 6-4 2015,approved me first – being water options 6-4 greatly keepapproved me well me

Model: Long-T5-TGlobal

Description: Long-document summarization with global attention

ROUGE Scores - R1: 0.3385 | R2: 0.0622 | RL: 0.1744

1234 OF 2019 STATE OF MAHARASHTRA ...APPELLANT VERSUS RAMESH KUMAR SHARMA ...RESPONDENT JUDGMENT This appeal arises from the judgment and order dated 15.03.2018 passed by the High Court of Bombay whereby the High Court acquitted the respondent of charges under Section 302 IPC (murder). Testimony of eyewitness who saw respondent leaving the deceased's residence at 8:45 PM The Trial Court convicted the respondent under Section 302 IPC and sentenced him to life imprisonment.

Model: Legal-LED

Description: Fine-tuned for legal and legislative text

ROUGE Scores - R1: 0.3772 | R2: 0.0973 | RL: 0.2105

On September 25, 2019, the Honorable A.K. Sikri of the Bombay High Court granted the SEC's request for an emergency asset freeze and other emergency relief. The SEC's complaint, filed in the U.S. District Court for the Central District of India, alleges that on 10.01.2015, Vijay Kumar was found dead in his residence with multiple stab wounds. According to the complaint, the respondent, who was the business partner of the deceased, was arrested based on circumstantial evidence including: Financial disputes between the deceased and respondent amounting to approximately Rs. 50 lakhs; Financial disputes between the deceased and respondent amounting to approximately Rs. 50 lakhs;

Model: T5-Large

Description: Text-to-text transfer transformer

ROUGE Scores - R1: 0.4308 | R2: 0.0829 | RL: 0.2154

a business partner of the deceased was found dead in his residence with multiple stab wounds . the trial court convicted the respondent under Section 302 IPC and sentenced him to life imprisonment . on appeal, the High Court reversed the conviction citing insufficient evidence and giving benefit of doubt to the accused . this appeal arises from the judgment and order

dated 15.03.2018 passed by the High Court of Bombay . acquittal was challenged by the appellant-State .

=====

=====

🏆 BEST PERFORMING MODEL: BART-Large-CNN

ROUGE-L Score: 0.2775

=====

=====

QNA:

📋 QUESTION-ANSWER PAIRS FOR EVALUATION:

Q1: What was the nature of the dispute between the deceased and the respondent?

Ground Truth: Financial disputes amounting to Rs. 50 lakhs

Type: factual

Q2: Under which section of IPC was the respondent originally convicted?

Ground Truth: Section 302 IPC

Type: factual

Q3: What was the final decision of the Supreme Court?

Ground Truth: The appeal was dismissed and the High Court's acquittal order was upheld

Type: reasoning

Q4: Which precedent case was cited by the Supreme Court?

Ground Truth: Sharad Birdhichand Sarda v. State of Maharashtra (1984) 4 SCC 116

Type: factual

Q5: What were the critical gaps in evidence identified by the Court?

Ground Truth: Uncertain time of death, contradictory eyewitness testimony, absence of blood stains on respondent's clothing, and broken chain of custody for murder weapon

Type: reasoning

=====

=====

🔍 LEGAL QUESTION ANSWERING - MODEL EVALUATION

=====

=====

Evaluating QA models: 100%

5/5 [01:19<00:00, 16.26s/it]

 Loading RoBERTa-SQuAD2...

Strong extractive QA baseline

config.json: 100%

571/571 [00:00<00:00, 56.5kB/s]

model.safetensors: 100%

496M/496M [00:06<00:00, 60.0MB/s]

Loading weights: 100%

199/199 [00:00<00:00, 695.83it/s, Materializing param=roberta.encoder.layer.11.output.dense.weight]

RobertaForQuestionAnswering LOAD REPORT from: deepset/roberta-base-squad2

Key	Status	
-----+-----+--+-		

roberta.embeddings.position_ids	UNEXPECTED	
---------------------------------	------------	--

Notes:

- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.

tokenizer_config.json: 100%

79.0/79.0 [00:00<00:00, 6.17kB/s]

vocab.json:

899k/? [00:00<00:00, 37.7MB/s]

merges.txt:

456k/? [00:00<00:00, 27.2MB/s]

special_tokens_map.json: 100%

772/772 [00:00<00:00, 88.8kB/s]

✓ Successfully loaded RoBERTa-SQuAD2

 Loading BERT-Large-SQuAD2...

Classical strong baseline for extractive QA

config.json: 100%

540/540 [00:00<00:00, 37.2kB/s]

model.safetensors: 100%

1.34G/1.34G [00:16<00:00, 129MB/s]

Loading weights: 100%

391/391 [00:00<00:00, 560.88it/s, Materializing param=qa_outputs.weight]

BertForQuestionAnswering LOAD REPORT from:

deepset/bert-large-uncased-whole-word-masking-squad2

Key	Status	
-----+-----+--+-		

bert.pooler.dense.weight	UNEXPECTED	
--------------------------	------------	--

bert.pooler.dense.bias	UNEXPECTED	
------------------------	------------	--

Notes:

- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.

tokenizer_config.json: 100%
19.0/19.0 [00:00<00:00, 1.52kB/s]
vocab.txt:
232k/? [00:00<00:00, 16.5MB/s]
added_tokens.json: 100%
2.00/2.00 [00:00<00:00, 237B/s]
special_tokens_map.json: 100%
112/112 [00:00<00:00, 13.6kB/s]
✓ Successfully loaded BERT-Large-SQuAD2

 Loading Legal-BERT...

Domain-adapted for legal information extraction

✗ Failed to load Legal-BERT: nsi319/legal-bert-base-uncased is not a local folder and is not a valid model identifier listed on '<https://huggingface.co/models>'
If this is a private repository, make sure to pass a token having permission to this repo either by logging in with `hf auth login` or by passing `token=<your\_token>`
Error with Legal-BERT: 'Legal-BERT'

 Loading Legal-BERT...

Domain-adapted for legal information extraction

✗ Failed to load Legal-BERT: nsi319/legal-bert-base-uncased is not a local folder and is not a valid model identifier listed on '<https://huggingface.co/models>'
If this is a private repository, make sure to pass a token having permission to this repo either by logging in with `hf auth login` or by passing `token=<your\_token>`
Error with Legal-BERT: 'Legal-BERT'

 Loading Legal-BERT...

Domain-adapted for legal information extraction

✗ Failed to load Legal-BERT: nsi319/legal-bert-base-uncased is not a local folder and is not a valid model identifier listed on '<https://huggingface.co/models>'
If this is a private repository, make sure to pass a token having permission to this repo either by logging in with `hf auth login` or by passing `token=<your\_token>`
Error with Legal-BERT: 'Legal-BERT'

 Loading Legal-BERT...

Domain-adapted for legal information extraction

✗ Failed to load Legal-BERT: nsi319/legal-bert-base-uncased is not a local folder and is not a valid model identifier listed on '<https://huggingface.co/models>'
If this is a private repository, make sure to pass a token having permission to this repo either by logging in with `hf auth login` or by passing `token=<your\_token>`
Error with Legal-BERT: 'Legal-BERT'

 Loading Legal-BERT...

Domain-adapted for legal information extraction

✗ Failed to load Legal-BERT: nsi319/legal-bert-base-uncased is not a local folder and is not a valid model identifier listed on '<https://huggingface.co/models>'
If this is a private repository, make sure to pass a token having permission to this repo either by logging in with `hf auth login` or by passing `token=<your\_token>`

Error with Legal-BERT: 'Legal-BERT'

⌚ Loading FLAN-T5-Base...

Instruction-tuned generative QA

config.json:

1.40k/? [00:00<00:00, 42.5kB/s]

tokenizer_config.json:

2.54k/? [00:00<00:00, 67.3kB/s]

spiece.model: 100%

792k/792k [00:00<00:00, 1.45MB/s]

tokenizer.json:

2.42M/? [00:00<00:00, 60.7MB/s]

special_tokens_map.json:

2.20k/? [00:00<00:00, 203kB/s]

model.safetensors: 100%

990M/990M [00:12<00:00, 45.6MB/s]

Loading weights: 100%

282/282 [00:00<00:00, 696.62it/s, Materializing param=shared.weight]

The tied weights mapping and config for this model specifies to tie shared.weight to lm_head.weight, but both are present in the checkpoints, so we will NOT tie them. You should update the config with `tie\_word\_embeddings=False` to silence this warning

generation_config.json: 100%

147/147 [00:00<00:00, 7.65kB/s]

✓ Successfully loaded FLAN-T5-Base

⌚ Loading DeBERTa-v3...

Improved attention disentanglement

config.json: 100%

579/579 [00:00<00:00, 65.0kB/s]

pytorch_model.bin: 100%

371M/371M [00:05<00:00, 81.2MB/s]

Loading weights: 100%

198/198 [00:00<00:00, 807.95it/s, Materializing param=deberta.encoder.rel_embeddings.weight]

DebertaV2ForQuestionAnswering LOAD REPORT from: microsoft/deberta-v3-base

Key	Status	
-----+-----+		

mask_predictions.LayerNorm.bias	UNEXPECTED	
---------------------------------	------------	--

mask_predictions.classifier.weight	UNEXPECTED	
------------------------------------	------------	--

lm_predictions.lm_head.LayerNorm.bias	UNEXPECTED	
---------------------------------------	------------	--

mask_predictions.LayerNorm.weight	UNEXPECTED	
-----------------------------------	------------	--

```

mask_predictions.classifier.bias      | UNEXPECTED |
lm_predictions.lm_head.bias          | UNEXPECTED |
lm_predictions.lm_head.LayerNorm.weight | UNEXPECTED |
lm_predictions.lm_head.dense.bias     | UNEXPECTED |
mask_predictions.dense.bias           | UNEXPECTED |
mask_predictions.dense.weight         | UNEXPECTED |
lm_predictions.lm_head.dense.weight   | UNEXPECTED |
qa_outputs.weight                    | MISSING   |
qa_outputs.bias                      | MISSING   |

```

Notes:

- UNEXPECTED :can be ignored when loading from different task/architecture; not ok if you expect identical arch.
- MISSING :those params were newly initialized because missing from the checkpoint. Consider training on your downstream task.

tokenizer_config.json: 100%

52.0/52.0 [00:00<00:00, 2.64kB/s]

spm.model: 100%

2.46M/2.46M [00:01<00:00, 2.29MB/s]

model.safetensors: 100%

371M/371M [00:03<00:00, 129MB/s]

✓ Successfully loaded DeBERTa-v3

```

=====
=====

```

QA MODEL PERFORMANCE COMPARISON

```

=====
=====

```

Model	Avg_F1	Avg_EM	Total_Questions
RoBERTa-SQuAD2	0.4044	0.0000	5
BERT-Large-SQuAD2	0.1533	0.0000	5
Legal-BERT	0.0000	0.0000	5
FLAN-T5-Base	0.2309	0.0000	5
DeBERTa-v3	0.0000	0.0000	5

```

=====
=====

```

DETAILED QA RESULTS BY MODEL

```

=====
=====

```

Model: RoBERTa-SQuAD2

Description: Strong extractive QA baseline
Average F1: 0.4044 | Average EM: 0.0000

Q1: What was the nature of the dispute between the deceased and the respondent?

Predicted: Financial disputes

Ground Truth: Financial disputes amounting to Rs. 50 lakhs

Q2: Under which section of IPC was the respondent originally convicted?

Predicted: Section 302

Ground Truth: Section 302 IPC

Q3: What was the final decision of the Supreme Court?

Predicted:

Ground Truth: The appeal was dismissed and the High Court's acquittal order was upheld

Q4: Which precedent case was cited by the Supreme Court?

Predicted: Sharad Birdhichand Sarda v. State of Maharashtra

Ground Truth: Sharad Birdhichand Sarda v. State of Maharashtra (1984) 4 SCC 116

Q5: What were the critical gaps in evidence identified by the Court?

Predicted:

Ground Truth: Uncertain time of death, contradictory eyewitness testimony, absence of blood stains on respondent's clothing, and broken chain of custody for murder weapon

Model: BERT-Large-SQuAD2

Description: Classical strong baseline for extractive QA

Average F1: 0.1533 | Average EM: 0.0000

Q1: What was the nature of the dispute between the deceased and the respondent?

Predicted: blood stains

Ground Truth: Financial disputes amounting to Rs. 50 lakhs

Q2: Under which section of IPC was the respondent originally convicted?

Predicted: 302

Ground Truth: Section 302 IPC

Q3: What was the final decision of the Supreme Court?

Predicted: acquitting the respondent is upheld

Ground Truth: The appeal was dismissed and the High Court's acquittal order was upheld

Q4: Which precedent case was cited by the Supreme Court?

Predicted:

Ground Truth: Sharad Birdhichand Sarda v. State of Maharashtra (1984) 4 SCC 116

Q5: What were the critical gaps in evidence identified by the Court?

Predicted:

Ground Truth: Uncertain time of death, contradictory eyewitness testimony, absence of blood stains on respondent's clothing, and broken chain of custody for murder weapon

Model: Legal-BERT

Description: Domain-adapted for legal information extraction

Average F1: 0.0000 | Average EM: 0.0000

Q1: What was the nature of the dispute between the deceased and the respondent?

Predicted: Error

Ground Truth: Financial disputes amounting to Rs. 50 lakhs

Q2: Under which section of IPC was the respondent originally convicted?

Predicted: Error

Ground Truth: Section 302 IPC

Q3: What was the final decision of the Supreme Court?

Predicted: Error

Ground Truth: The appeal was dismissed and the High Court's acquittal order was upheld

Q4: Which precedent case was cited by the Supreme Court?

Predicted: Error

Ground Truth: Sharad Birdhichand Sarda v. State of Maharashtra (1984) 4 SCC 116

Q5: What were the critical gaps in evidence identified by the Court?

Predicted: Error

Ground Truth: Uncertain time of death, contradictory eyewitness testimony, absence of blood stains on respondent's clothing, and broken chain of custody for murder weapon

Model: FLAN-T5-Base

Description: Instruction-tuned generative QA

Average F1: 0.2309 | Average EM: 0.0000

Q1: What was the nature of the dispute between the deceased and the respondent?

Predicted: Financial disputes between the deceased and the respondent

Ground Truth: Financial disputes amounting to Rs. 50 lakhs

Q2: Under which section of IPC was the respondent originally convicted?

Predicted: Murder

Ground Truth: Section 302 IPC

Q3: What was the final decision of the Supreme Court?

Predicted: High Court acquitted the respondent of charges

Ground Truth: The appeal was dismissed and the High Court's acquittal order was upheld

Q4: Which precedent case was cited by the Supreme Court?

Predicted: Sharad Birdhichand Sarda

Ground Truth: Sharad Birdhichand Sarda v. State of Maharashtra (1984) 4 SCC 116

Q5: What were the critical gaps in evidence identified by the Court?

Predicted: 1. The exact time of death could not be established beyond reasonable doubt 2.

The eyewitness testimony was contradictory regarding the time

Ground Truth: Uncertain time of death, contradictory eyewitness testimony, absence of blood stains on respondent's clothing, and broken chain of custody for murder weapon

Model: DeBERTa-v3

Description: Improved attention disentanglement

Average F1: 0.0000 | Average EM: 0.0000

Q1: What was the nature of the dispute between the deceased and the respondent?

Predicted: Sikri)

(Justice

Ground Truth: Financial disputes amounting to Rs. 50 lakhs

Q2: Under which section of IPC was the respondent originally convicted?

Predicted: Sikri)

(Justice

Ground Truth: Section 302 IPC

Q3: What was the final decision of the Supreme Court?

Predicted: Sikri)

(Justice

Ground Truth: The appeal was dismissed and the High Court's acquittal order was upheld

Q4: Which precedent case was cited by the Supreme Court?

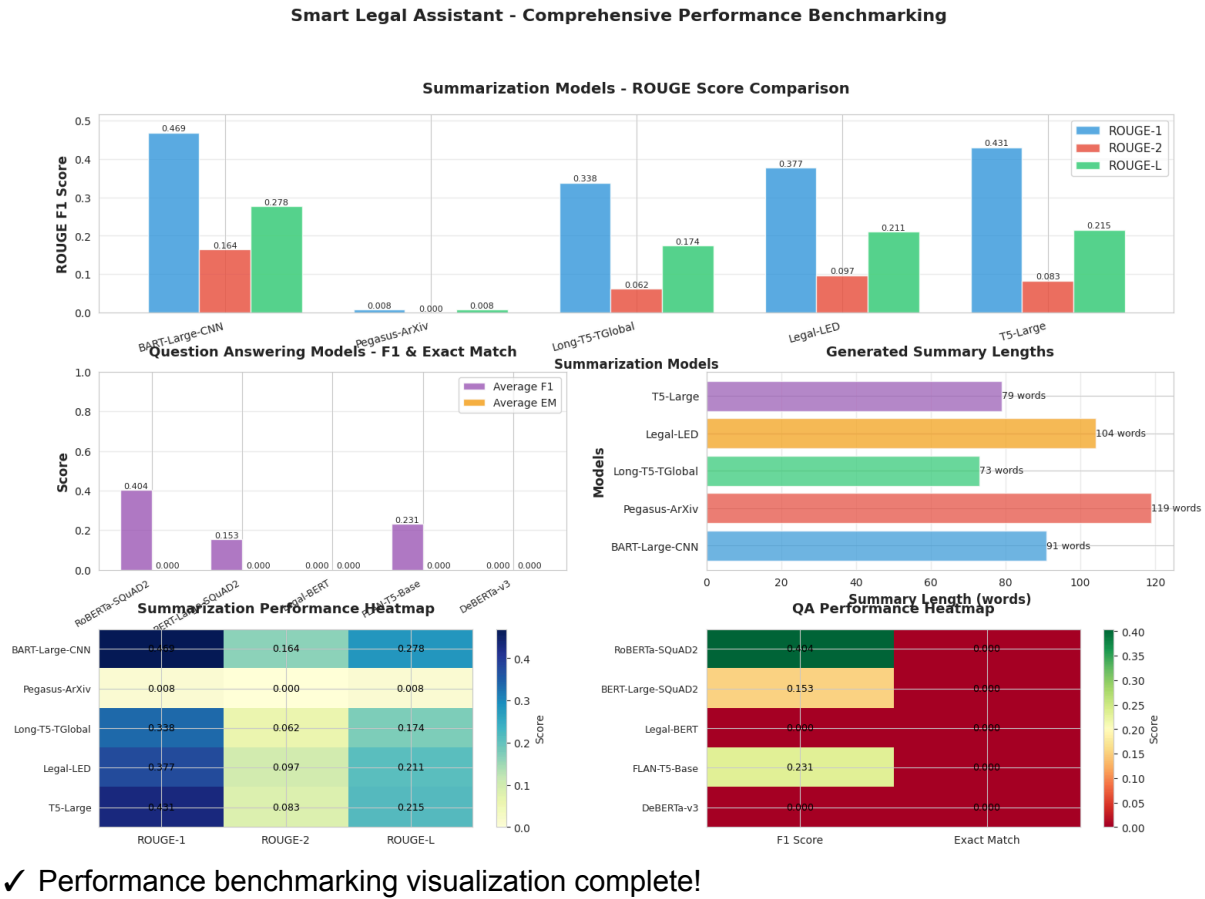
Predicted: Sikri)

(Justice

Ground Truth: Sharad Birdhichand Sarda v. State of Maharashtra (1984) 4 SCC 116

Q5: What were the critical gaps in evidence identified by the Court?
Predicted: Sikri)
(Justice
Ground Truth: Uncertain time of death, contradictory eyewitness testimony, absence of blood stains on respondent's clothing, and broken chain of custody for murder weapon

🏆 BEST PERFORMING QA MODEL: RoBERTa-SQuAD2
Average F1 Score: 0.4044



✓ Performance benchmarking visualization complete!

📊 COMPREHENSIVE PERFORMANCE REPORT

1. SUMMARIZATION MODELS - DETAILED METRICS

Model	Description	ROUGE-1	ROUGE-2	ROUGE-L	
Summary_Length					
BART-Large-CNN	State-of-the-art abstractive summarizer	0.4689	0.1643	0.2775	91
Pegasus-ArXiv Gap	Sentence Generation for superior ROUGE scores	0.0081	0.0000		119
0.0081					
Long-T5-TGlobal	Long-document summarization with global attention	0.3385	0.0622		73
0.1744					
Legal-LED	Fine-tuned for legal and legislative text	0.3772	0.0973	0.2105	104
T5-Large	Text-to-text transfer transformer	0.4308	0.0829	0.2154	79

Statistical Summary (Summarization):

Best ROUGE-1: BART-Large-CNN (0.4689)
Best ROUGE-2: BART-Large-CNN (0.1643)
Best ROUGE-L: BART-Large-CNN (0.2775)
Average ROUGE-L across all models: 0.1772

2. QUESTION ANSWERING MODELS - DETAILED METRICS

Model	Description	Avg_F1	Avg_EM	Total_Questions	
RoBERTa-SQuAD2	Strong extractive QA baseline	0.4044	0.0000		5
BERT-Large-SQuAD2	Classical strong baseline for extractive QA	0.1533	0.0000		5
Legal-BERT	Domain-adapted for legal information extraction	0.0000	0.0000		5
FLAN-T5-Base	Instruction-tuned generative QA	0.2309	0.0000		5
DeBERTa-v3	Improved attention disentanglement	0.0000	0.0000		5

Statistical Summary (Question Answering):

Best F1 Score: RoBERTa-SQuAD2 (0.4044)
Best Exact Match: RoBERTa-SQuAD2 (0.0000)
Average F1 across all models: 0.1577
Average EM across all models: 0.0000

3. MODEL RECOMMENDATIONS

🏆 RECOMMENDED SUMMARIZATION MODEL: BART-Large-CNN
Rationale: Highest ROUGE-L score (0.2775)
State-of-the-art abstractive summarizer

🏆 RECOMMENDED QA MODEL: RoBERTa-SQuAD2
Rationale: Highest F1 score (0.4044)
Strong extractive QA baseline

=====
=====

✓ Performance report generation complete!

=====
=====